
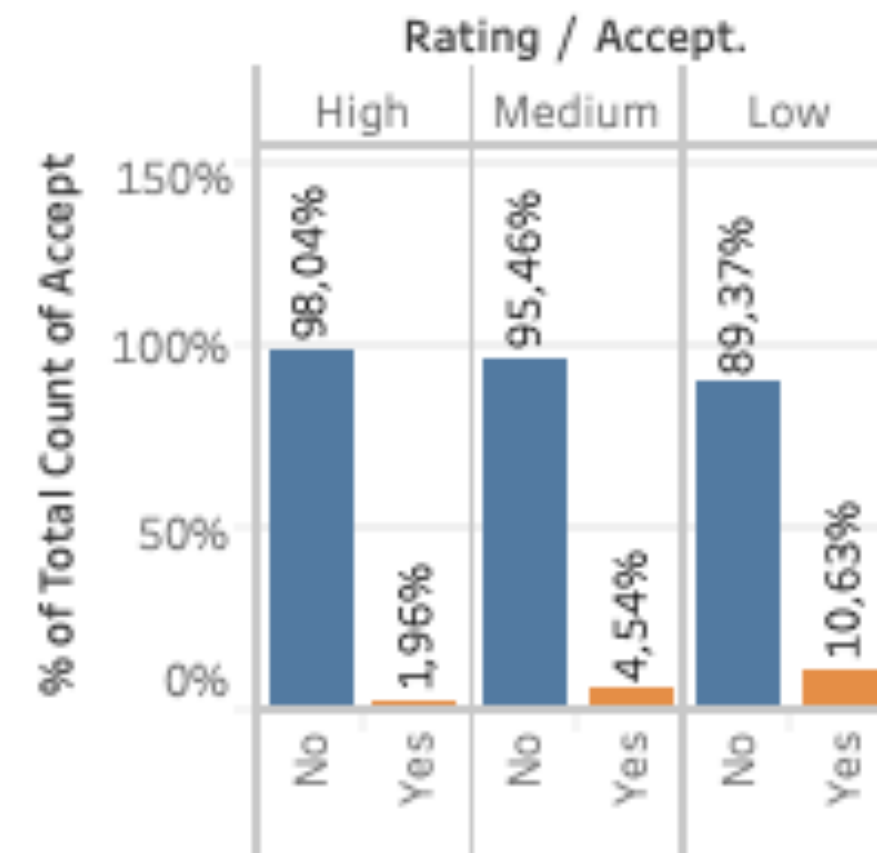






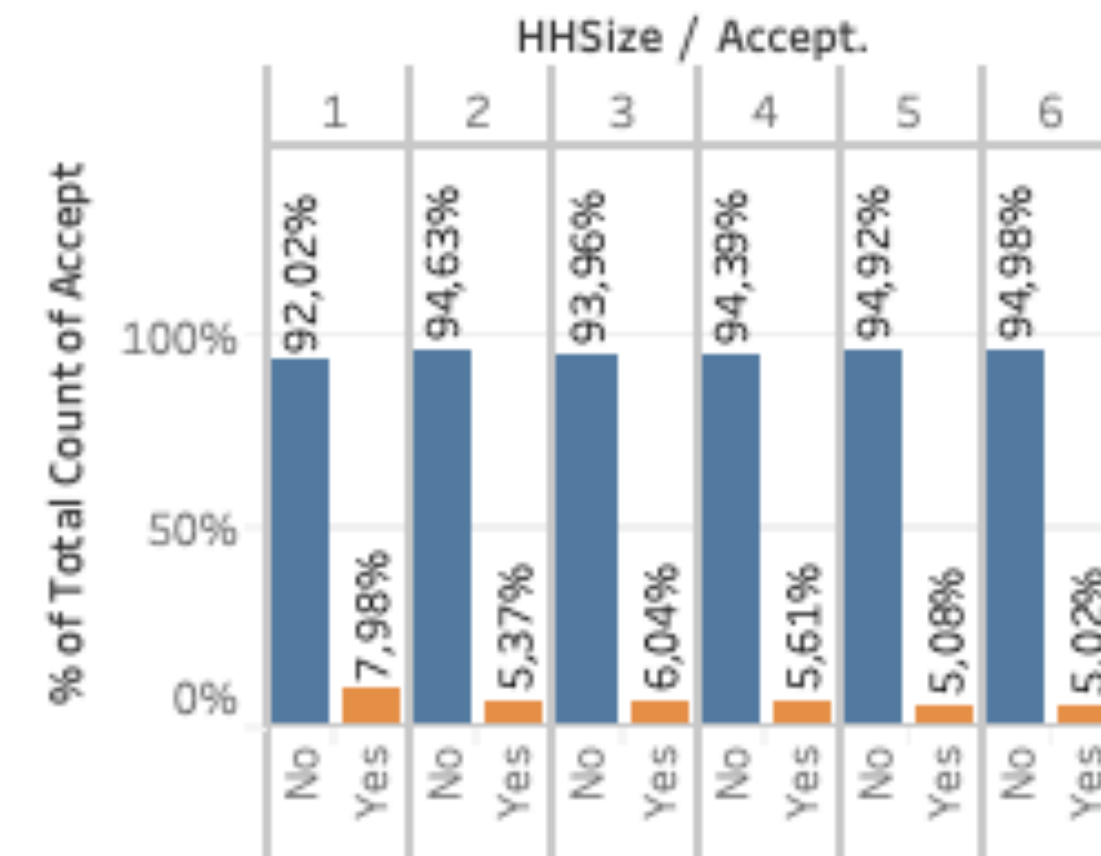
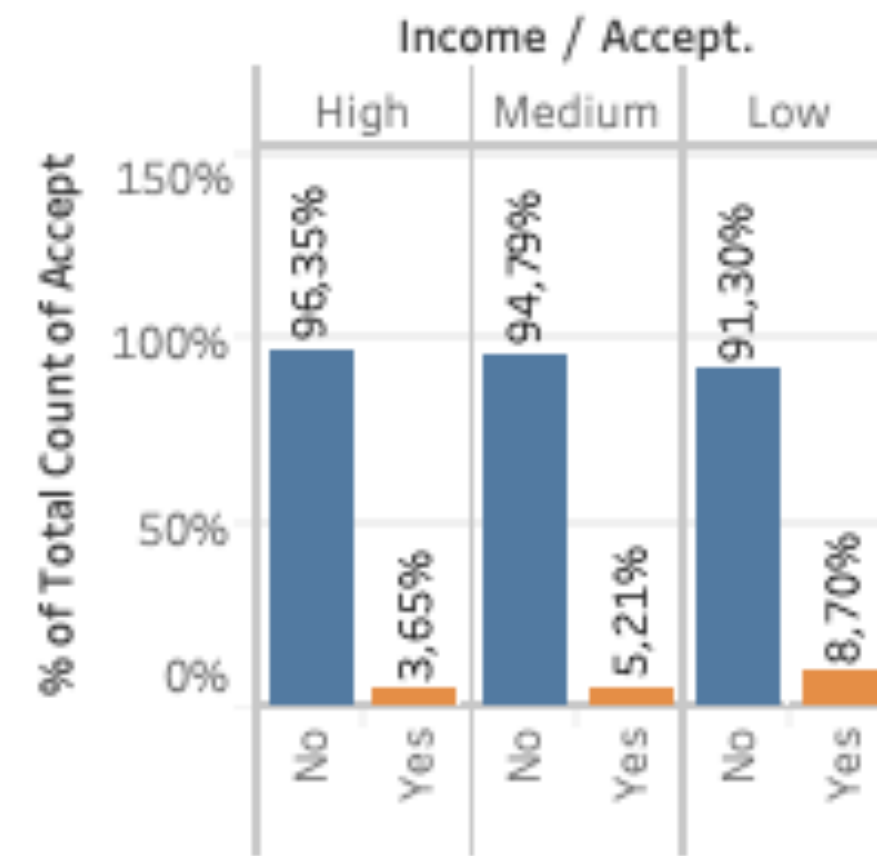



---



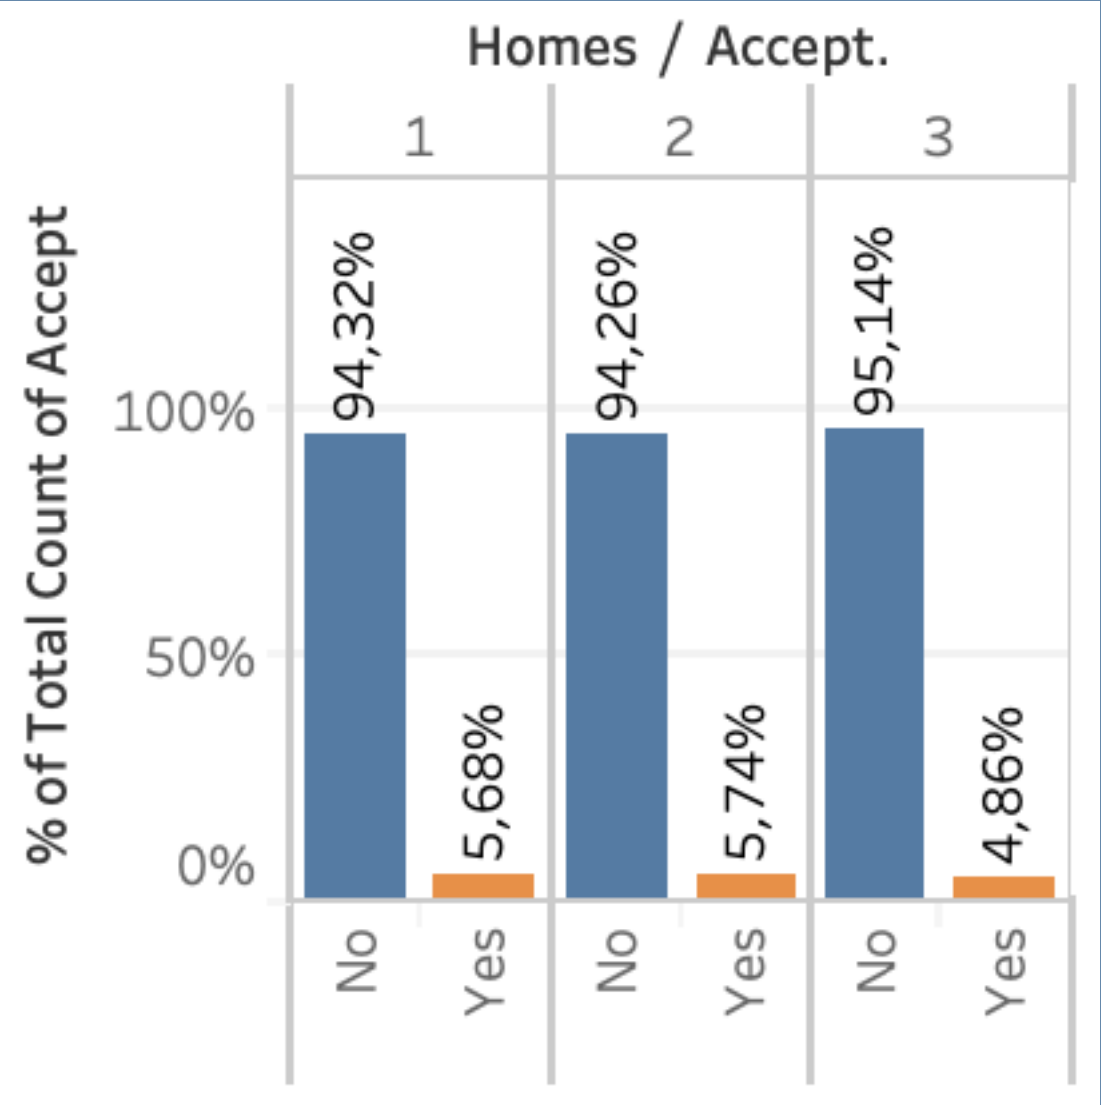


Will the customers we contact accept a Credit Card offer?

# CLASSIFICATION MODEL

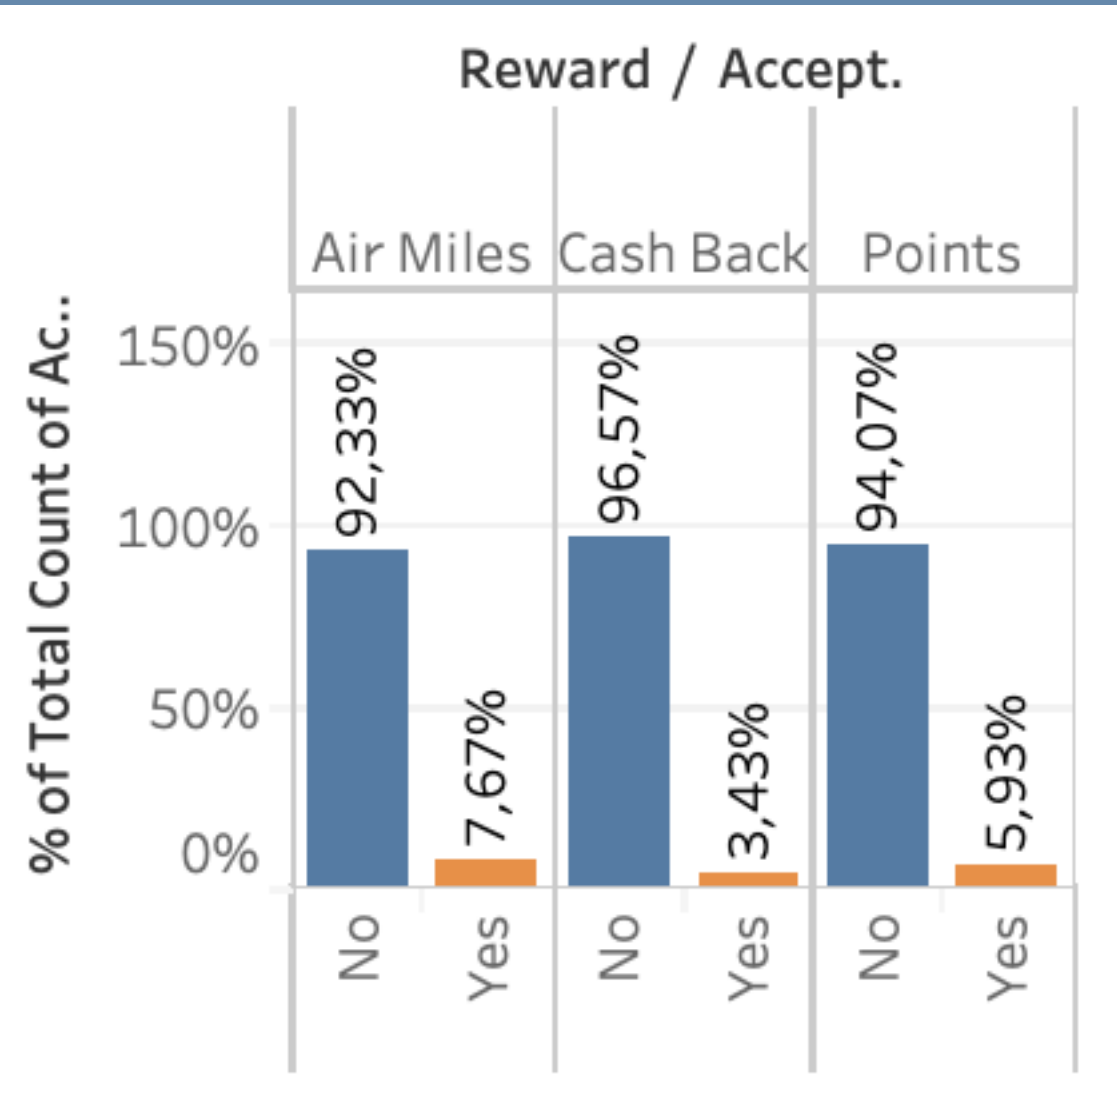
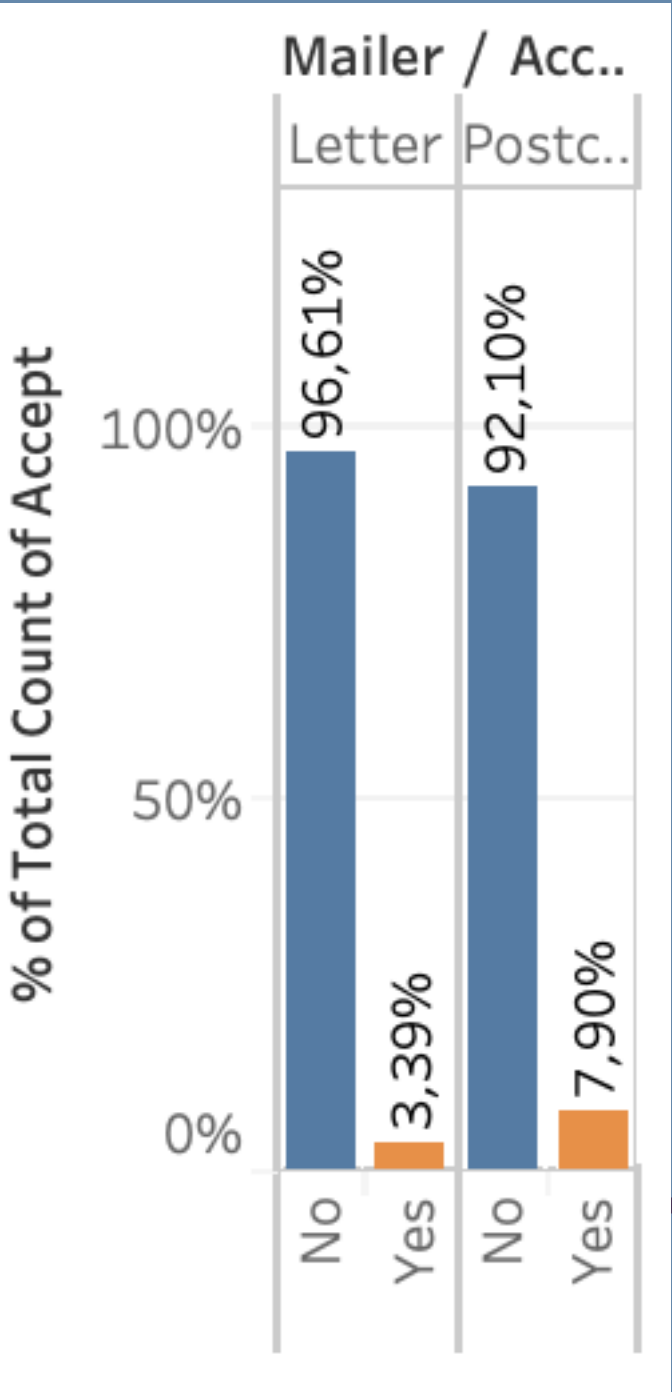


# FAST FIRST DATA VIZ



SEEMINGLY NOT  
SIGNIFICANT  
TRAITS



POTENTIALLY  
SIGNIFICANT  
TRAITS







## OTHER TRAITS THAT COULD BE IMPORTANT

```
data['LiveAlone'] = np.where(data['HHSize'] == 1, 1, 0)
```

```
data['LowBalance'] = np.where(data['AvgBal'] <= 500, 1, 0)  
data['HighBalance'] = np.where(data['AvgBal'] >= 1000, 1, 0)
```



# DATASET

Accept	AirMiles	Points	Letter	Income	Rating	LiveAlone	LowBalance	AvgBal
0	0	1	0	1	3	0	0	1175
0	0	1	1	2	1	0	0	811
0	0	0	1	3	3	0	0	1754
0	1	0	0	3	3	0	0	689
0	0	0	1	1	1	0	0	1018
0	0	0	0	2	3	0	0	1130
0	1	0	0	1	1	1	0	877
1	1	0	1	2	1	0	0	769
0	0	0	0	2	1	0	0	709
0	0	1	0	3	2	1	0	690
0	0	1	1	2	3	0	0	863
0	0	0	0	1	1	0	0	733
0	0	1	0	1	3	0	0	771
0	1	0	0	1	2	0	0	659
0	1	0	1	2	1	0	0	1258

---

# MAIN ISSUE: DATA IMBALANCE

5.68%      94.32%

YES

NO

**NORMAL MODEL WILL BE BIASED TOWARDS  
PREDICTING NO FOR ALL OBSERVATIONS**

---

---

# **TREATING DATA IMBALANCE**

**NORMAL MODEL WILL BE BIASED  
TOWARDS PREDICTING NO FOR ALL  
OBSERVATIONS**

**OUR JOB IS TO MAKE A MODEL THAT  
FOCUSES ON PREDICTING YESES  
AS ACCURATELY AS POSSIBLE.**

---





# 1.A LOGISTIC MODEL WITH COMPLETE DATASET

```
1 # This is the model for the complete dataset, df1.
2
3 for sam in [1, 0.9, 0.8, 0.66, 0.5, 0.4, 0.33, 0.25, 0.2]:
4     for test in [0.2, 0.25, 0.3]:
5         undersample = RandomUnderSampler(sampling_strategy=sam)
6         X_und, y_und = undersample.fit_resample(X1, y1)
7         X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
8         classifier = LogisticRegression(solver='liblinear', class_weight='balanced')
9         classifier.fit(X_train, y_train)
10        y_pred = classifier.predict(X_test)
11        pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
12        acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
13        print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}: {acc:.2f}% of sensitivity and a
```

```
Log model with a 1.00 ratio, test size of 0.20: 73.50% of sensitivity and a 63.64% of precision.
Log model with a 1.00 ratio, test size of 0.25: 75.21% of sensitivity and a 67.66% of precision.
Log model with a 1.00 ratio, test size of 0.30: 75.51% of sensitivity and a 70.70% of precision.
Log model with a 0.90 ratio, test size of 0.20: 64.79% of sensitivity and a 66.99% of precision.
Log model with a 0.90 ratio, test size of 0.25: 66.54% of sensitivity and a 64.29% of precision.
Log model with a 0.90 ratio, test size of 0.30: 67.54% of sensitivity and a 67.54% of precision.
```



# 1.B LOGISTIC MODEL WITH REDUCED DATASET

```
1 # This is the model for the reduced dataset, df2.
2
3 for sam in [1, 0.9, 0.8, 0.66, 0.5, 0.4, 0.33, 0.25, 0.2]:
4     for test in [0.2, 0.25, 0.3]:
5         undersample = RandomUnderSampler(sampling_strategy=sam)
6         X_und, y_und = undersample.fit_resample(X2, y2)
7         X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
8         classifier = LogisticRegression(solver='liblinear', class_weight='balanced')
9         classifier.fit(X_train, y_train)
10        y_pred = classifier.predict(X_test)
11        pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
12        acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
13        print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}: {acc:.2f}% of sensitivity and a
```

Log model with a 1.00 ratio, test size of 0.20: 74.50% of sensitivity and a 69.63% of precision.

Log model with a 1.00 ratio, test size of 0.25: 81.40% of sensitivity and a 68.17% of precision.

Log model with a 1.00 ratio, test size of 0.30: 76.19% of sensitivity and a 66.87% of precision.

Log model with a 0.90 ratio, test size of 0.20: 66.67% of sensitivity and a 68.27% of precision.

Log model with a 0.90 ratio, test size of 0.25: 67.32% of sensitivity and a 64.07% of precision.

Log model with a 0.90 ratio, test size of 0.30: 71.15% of sensitivity and a 63.82% of precision.

Log model with a 0.80 ratio, test size of 0.20: 74.26% of sensitivity and a 64.38% of precision.

## 2. K-NEAREST NEIGHBOURS MODEL

```
1 # This is the model for the reduced dataset, df2.
2
3 for sam in [1, 0.9, 0.8]:
4     for test in [0.2, 0.25, 0.3]:
5         for k in [3, 5, 7, 9]:
6             undersample = RandomUnderSampler(sampling_strategy=sam)
7             X_und, y_und = undersample.fit_resample(X2, y2)
8             X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
9             nbrs = NearestNeighbors(n_neighbors=k)
10            nbrs.fit(X_train, y_train)
11            y_pred = classifier.predict(X_test)
12            pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
13            acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
14            print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}, K = {k:.0f}: {acc:.2f}% of sens.
```

```
Log model with a 1.00 ratio, test size of 0.20, K = 3: 72.50% of sensitivity and a 68.08% of precision.
Log model with a 1.00 ratio, test size of 0.20, K = 5: 72.50% of sensitivity and a 68.72% of precision.
Log model with a 1.00 ratio, test size of 0.20, K = 7: 72.50% of sensitivity and a 68.72% of precision.
Log model with a 1.00 ratio, test size of 0.20, K = 9: 72.50% of sensitivity and a 70.73% of precision.
```



### 3. SUPPORT VECTOR MACHINE (SVM)

```
1 # This is the model for the reduced dataset, df2.
2
3 for sam in [1, 0.9, 0.8]:
4     for test in [0.2, 0.25, 0.3]:
5         undersample = RandomUnderSampler(sampling_strategy=sam)
6         X_und, y_und = undersample.fit_resample(X2, y2)
7         X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
8         clas = svm.SVC()
9         clas.fit(X_train, y_train)
10        y_pred = classifier.predict(X_test)
11        pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
12        acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
13        print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}, K = {k:.0f}: {acc:.2f}% of sens:
```

```
Log model with a 1.00 ratio, test size of 0.20, K = 9: 72.50% of sensitivity and a 65.91% of precision.
Log model with a 1.00 ratio, test size of 0.25, K = 9: 73.55% of sensitivity and a 70.08% of precision.
Log model with a 1.00 ratio, test size of 0.30, K = 9: 74.49% of sensitivity and a 68.87% of precision.
Log model with a 0.90 ratio, test size of 0.20, K = 9: 64.79% of sensitivity and a 68.32% of precision.
```

## 4. DECISION TREE MODEL

```
1 # This is the model for the reduced dataset, df2.
2
3 for sam in [1, 0.9, 0.8]:
4     for test in [0.2, 0.25, 0.3]:
5         undersample = RandomUnderSampler(sampling_strategy=sam)
6         X_und, y_und = undersample.fit_resample(X2, y2)
7         X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
8         clas = tree.DecisionTreeClassifier()
9         clas.fit(X_train, y_train)
10        y_pred = classifier.predict(X_test)
11        pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
12        acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
13        print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}, K = {k:.0f}: {acc:.2f}% of sensi
```

```
Log model with a 1.00 ratio, test size of 0.20, K = 9: 72.50% of sensitivity and a 73.98% of precision.
Log model with a 1.00 ratio, test size of 0.25, K = 9: 73.55% of sensitivity and a 66.92% of precision.
Log model with a 1.00 ratio, test size of 0.30, K = 9: 74.49% of sensitivity and a 67.18% of precision.
Log model with a 0.90 ratio, test size of 0.20, K = 9: 64.79% of sensitivity and a 67.65% of precision.
```



# 5. RANDOM FOREST MODEL

```
1 # This is the model for the reduced dataset, df2.
2
3 for sam in [1, 0.9, 0.8]:
4     for test in [0.2, 0.25, 0.3]:
5         undersample = RandomUnderSampler(sampling_strategy=sam)
6         X_und, y_und = undersample.fit_resample(X2, y2)
7         X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=test, random_state=42)
8         clasRF = RandomForestClassifier(max_depth=5, random_state=42, n_estimators = 100)
9         clasRF.fit(X_train, y_train)
10        y_pred = classifier.predict(X_test)
11        pre = classification_report(y_test, y_pred, output_dict=True)['1']['precision']*100
12        acc = classification_report(y_test, y_pred, output_dict=True)['1']['recall']*100
13        print(f'Log model with a {sam:.2f} ratio, test size of {test:.2f}, K = {k:.0f}: {acc:.2f}% of sensi
```

```
Log model with a 1.00 ratio, test size of 0.20, K = 9: 72.50% of sensitivity and a 69.05% of precision.
Log model with a 1.00 ratio, test size of 0.25, K = 9: 73.55% of sensitivity and a 69.80% of precision.
Log model with a 1.00 ratio, test size of 0.30, K = 9: 74.49% of sensitivity and a 72.04% of precision.
Log model with a 0.90 ratio, test size of 0.20, K = 9: 64.79% of sensitivity and a 62.73% of precision.
Log model with a 0.80 ratio, test size of 0.20, K = 9: 66.83% of sensitivity and a 66.83% of precision.
```

# CHOSEN MODEL STATISTICS

```
1 undersample = RandomUnderSampler(sampling_strategy=1.0)
2 X_und, y_und = undersample.fit_resample(X2, y2)
3 X_train, X_test, y_train, y_test = train_test_split(X_und, y_und, test_size=0.25, random_state=42)
4 classifier = LogisticRegression(solver='liblinear', class_weight='balanced')
5 classifier.fit(X_train, y_train)
6 y_pred = classifier.predict(X_test)
7 print(classification_report(y_test, y_pred, output_dict=False))
```

	precision	recall	f1-score	support
0	0.79	0.64	0.70	269
1	0.67	0.81	0.73	242
accuracy			0.72	511
macro avg	0.73	0.72	0.72	511
weighted avg	0.73	0.72	0.72	511

